

Big data: Classification (Titanic dataset)

Ignacio Amaya and Justin Vailhere

November 4, 2015

Getting and cleaning the data

We read the Titanic file from the subject folder.

```
titanic <- read.csv("C:\\Users\\Ignacio\\Documents\\MIS DOCUMENTOS\\DATA SCIENCE MASTER (EIT DIGITAL)\\")
```

The variables in this data set are:

Key	Description
pclass	Passenger Class (1 = 1st; 2 = 2nd; 3 = 3rd)
survival	Survival (0 = No; 1 = Yes)
name	Name
sex	Sex
age	Age
sibsp	Number of Siblings/Spouses Aboard
parch	Number of Parents/Children Aboard
ticket	Ticket Number
fare	Passenger Fare
cabin	Cabin
embarked	Port of Embarkation (C,Q or S)
boat	Lifeboat
body	Body Identification Number
home.dest	Home/Destination

We can have an overview of all the values we have.

```
summary(titanic)
```

```
##   PassengerId      Survived  Pclass
##   Min.   : 1.0    Min.   :0.0000   Min.   :1.000
##   1st Qu.:223.5    1st Qu.:0.0000   1st Qu.:2.000
##   Median :446.0    Median :0.0000   Median :3.000
##   Mean   :446.0    Mean   :0.3838   Mean   :2.309
##   3rd Qu.:668.5    3rd Qu.:1.0000   3rd Qu.:3.000
##   Max.   :891.0    Max.   :1.0000   Max.   :3.000
##
##                                     Name      Sex      Age
##   Abbing, Mr. Anthony                : 1   female:314   Min.   : 0.42
##   Abbott, Mr. Rossmore Edward         : 1   male  :577   1st Qu.:20.12
##   Abbott, Mrs. Stanton (Rosa Hunt)    : 1                                Median :28.00
##   Abelson, Mr. Samuel                 : 1                                Mean   :29.70
##   Abelson, Mrs. Samuel (Hannah Wizo): 1                                3rd Qu.:38.00
##   Adahl, Mr. Mauritz Nils Martin       : 1                                Max.   :80.00
##   (Other)                             :885                                NA's   :177
```

```
##      SibSp      Parch      Ticket      Fare
## Min.   :0.000 Min.   :0.0000 1601    : 7 Min.   : 0.00
## 1st Qu.:0.000 1st Qu.:0.0000 347082 : 7 1st Qu.: 7.91
## Median :0.000 Median :0.0000 CA. 2343: 7 Median : 14.45
## Mean   :0.523 Mean   :0.3816 3101295 : 6 Mean   : 32.20
## 3rd Qu.:1.000 3rd Qu.:0.0000 347088 : 6 3rd Qu.: 31.00
## Max.   :8.000 Max.   :6.0000 CA 2144 : 6 Max.   :512.33
##                                     (Other) :852
##      Cabin      Embarked
## B96 B98      : 4 C      :168
## C23 C25 C27: 4 Q      : 77
## G6           : 4 S      :644
## C22 C26      : 3 NA's: 2
## D           : 3
## (Other)      :186
## NA's         :687
```

Some variables stored in the dataframe are treated as numerical, but they are categorical. So we are going to transform them into categorical variables.

```
titanic$Survived=as.factor(titanic$Survived)
titanic$Pclass=as.factor(titanic$Pclass)
titanic$SibSp=as.factor(titanic$SibSp)
titanic$Parch=as.factor(titanic$Parch)
```

Question 1: Which are the variables that are discarded and why? We want to discard those variables that don't give us any information in order to predict if a passenger lives or dies. There are some variables that don't give us any extra information, because they have a lot of missing values (for example, the cabin variable). Others are also discarded because their values are all different (for example, the names), so we can't extract information from them.

```
titanic_s1=subset.data.frame(titanic,select=c(Survived,Pclass,Sex,Age,SibSp,Parch,Fare,Embarked))
```

We can get a summary with the information of the subset we've selected.

```
summary(titanic_s1)
```

```
## Survived Pclass      Sex      Age      SibSp      Parch
## 0:549      1:216 female:314 Min.   : 0.42 0:608 0:678
## 1:342      2:184 male  :577 1st Qu.:20.12 1:209 1:118
##           3:491      Median :28.00 2: 28 2: 80
##           Mean   :29.70 3: 16 3: 5
##           3rd Qu.:38.00 4: 18 4: 4
##           Max.   :80.00 5: 5 5: 5
##           NA's   :177 8: 7 6: 1
##      Fare      Embarked
## Min.   : 0.00 C      :168
## 1st Qu.: 7.91 Q      : 77
## Median : 14.45 S      :644
## Mean   : 32.20 NA's: 2
## 3rd Qu.: 31.00
## Max.   :512.33
##
```

We observe that we have 177 NA in age and 2 NA in Embarked. We don't want to have missing values, so we remove the entries where those NA appear. Our final cleaned dataset has 712 observations of 8 variables.

```
titanic_cleaned=na.omit(titanic_s1,titanic_s1$Age,titanic_s1$Embarked)
summary(titanic_cleaned)
```

```
##   Survived Pclass      Sex      Age      SibSp  Parch
##   0:424    1:184  female:259  Min.   : 0.42  0:469    0:519
##   1:288    2:173   male  :453  1st Qu.:20.00 1:183    1:110
##           3:355           Median :28.00 2: 25    2: 68
##           Mean   :29.64 3: 12    3:  5
##           3rd Qu.:38.00 4: 18    4:  4
##           Max.   :80.00 5:  5    5:  5
##                        8:  0    6:  1
##      Fare      Embarked
##   Min.   : 0.00  C:130
##   1st Qu.: 8.05  Q: 28
##   Median :15.65  S:554
##   Mean   :34.57
##   3rd Qu.:33.00
##   Max.   :512.33
##
```

Training a single decision tree

Question 2: Plot the tree with the `prp()` function (package `rpart.plot` that we must install if it is not already installed). Which are the most important variables according to the tree? Does the tree change if we rerun the `rpart()` function? Why? First of all, we load the packages we need.

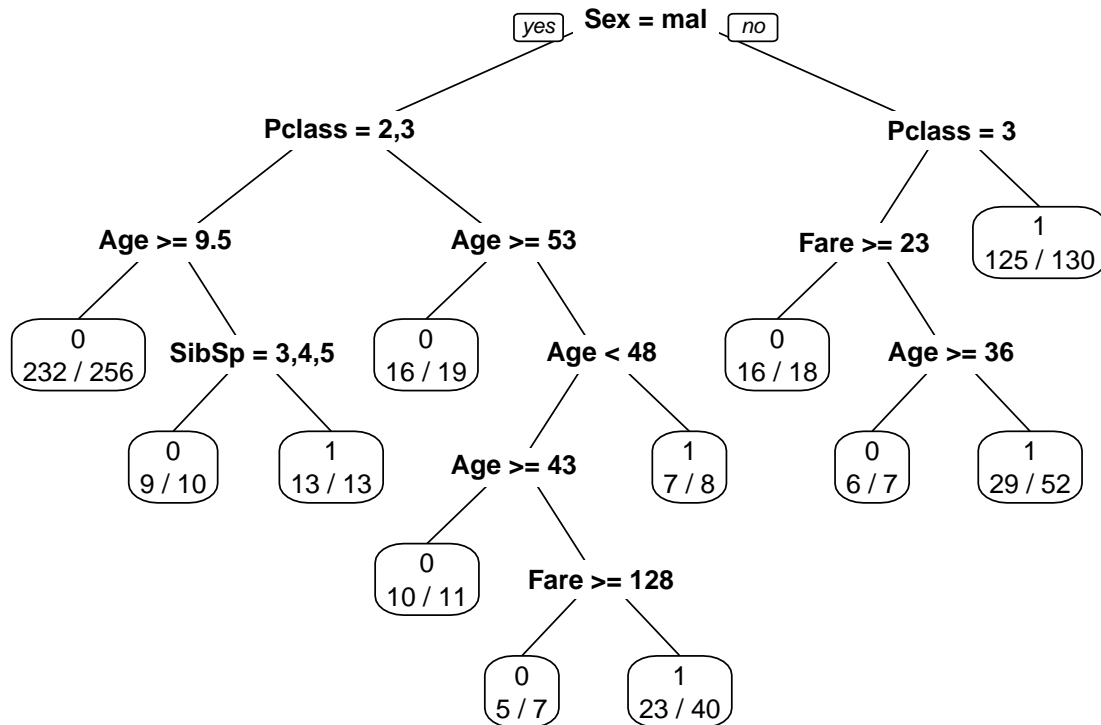
```
library(caret)
library(rpart)
library(rpart.plot)
```

We set the seed, so that way all the values we are going to get each time we execute the code are going to be the same. We divide our data into a training set and a testing set (80% of the entries are going to the training set and 20% to the testing set).

```
set.seed(100)
partition=createDataPartition(titanic_cleaned$Survived,p = 0.8,list=FALSE)
train_set=titanic_cleaned[partition,]
test_set=titanic_cleaned[-partition,]

random_tree=rpart(train_set)

prp(random_tree,extra=2)
```

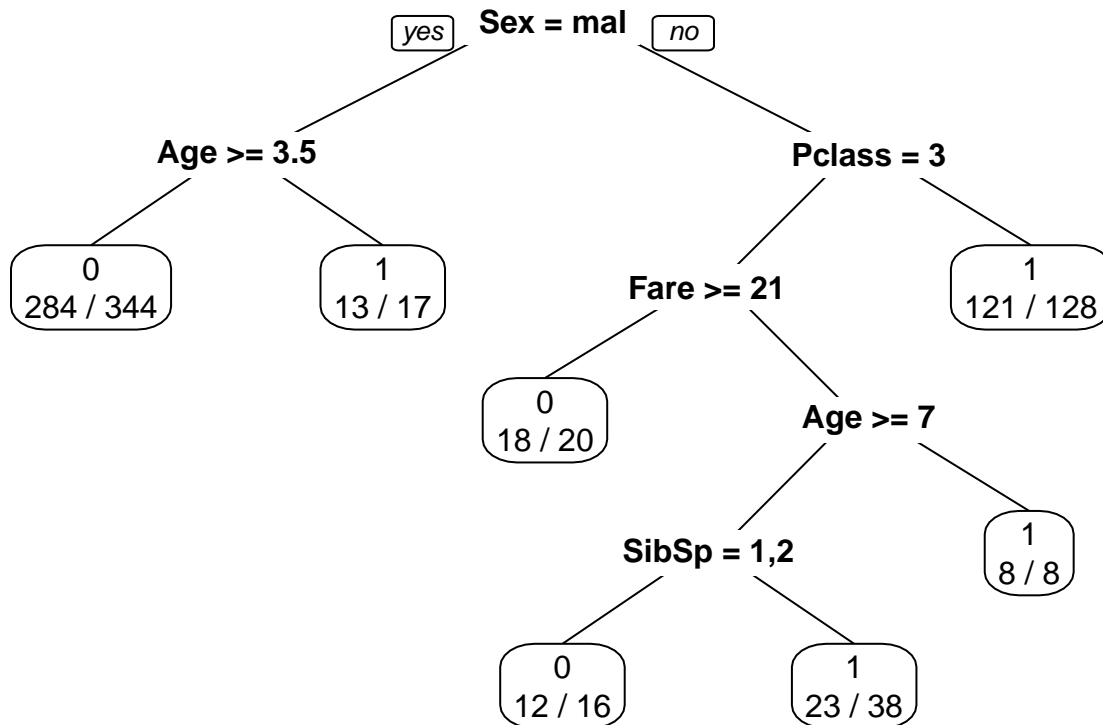


The most important variables are the sex and the Pclass. If the passenger is a woman and is not in third class, then the probability of her surviving is very high (only 5 out of 130 in this case died). If we rerun only the rpart function the tree obtained doesn't change. This is because the partition used is the same. But if we rerun the partition without setting the seed we would obtain different results, as we have a different partition. An example of this would be:

```

partition2=createDataPartition(titanic_cleaned$Survived,p = 0.8,list=FALSE)
train_set2=titanic_cleaned[partition2,]
test_set2=titanic_cleaned[-partition2,]
random_tree2=rpart(train_set2)
prp(random_tree2,extra=2)

```



```
testPred <- predict(random_tree, test_set,type="class")
confusionMatrix(testPred,test_set$Survived)
```

Question 3: Which are the performance values (accuracy, sensitivity, specificity, etc.) of the learned model on the testing subset?

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0  1
##           0 63 14
##           1 21 43
##
##           Accuracy : 0.7518
##           95% CI : (0.6721, 0.8206)
##           No Information Rate : 0.5957
##           P-Value [Acc > NIR] : 7.372e-05
##
##           Kappa : 0.4946
##           McNemar's Test P-Value : 0.3105
##
##           Sensitivity : 0.7500
##           Specificity : 0.7544
```

```
##          Pos Pred Value : 0.8182
##          Neg Pred Value : 0.6719
##          Prevalence : 0.5957
##          Detection Rate : 0.4468
##          Detection Prevalence : 0.5461
##          Balanced Accuracy : 0.7522
##
##          'Positive' Class : 0
##
```

As we can see the performance is good. We have an accuracy above 75% and the sensitivity (true positive rate) and specificity (true negative rate) are also good. As they have similar values we know that the predictions of people dying and surviving have the same rate of success. In the confusion we can see that the decision tree failed in 35 cases out of 141 (number of entries in our testing set).

Question 4: We have learned a decision tree model with the CART algorithm. Which are the values for the parameters of this algorithm (minbucket, minsplit, complexity parameter, cost, etc.)? (CLUE: check the `rpart()` function documentation.) The default parameters for the `rpart` function are: `rpart.control(minsplit = 20, minbucket = round(minsplit/3), cp = 0.01, maxcompete = 4, maxsurrogate = 5, usesurrogate = 2, xval = 10, surrogatestyle = 0, maxdepth = 30, ...)`

cost

A vector of non-negative costs, one for each variable in the model. Defaults to one for all variables. These are scalings to be applied when considering splits, so the improvement on splitting on a variable is divided by its cost in deciding which split to choose.

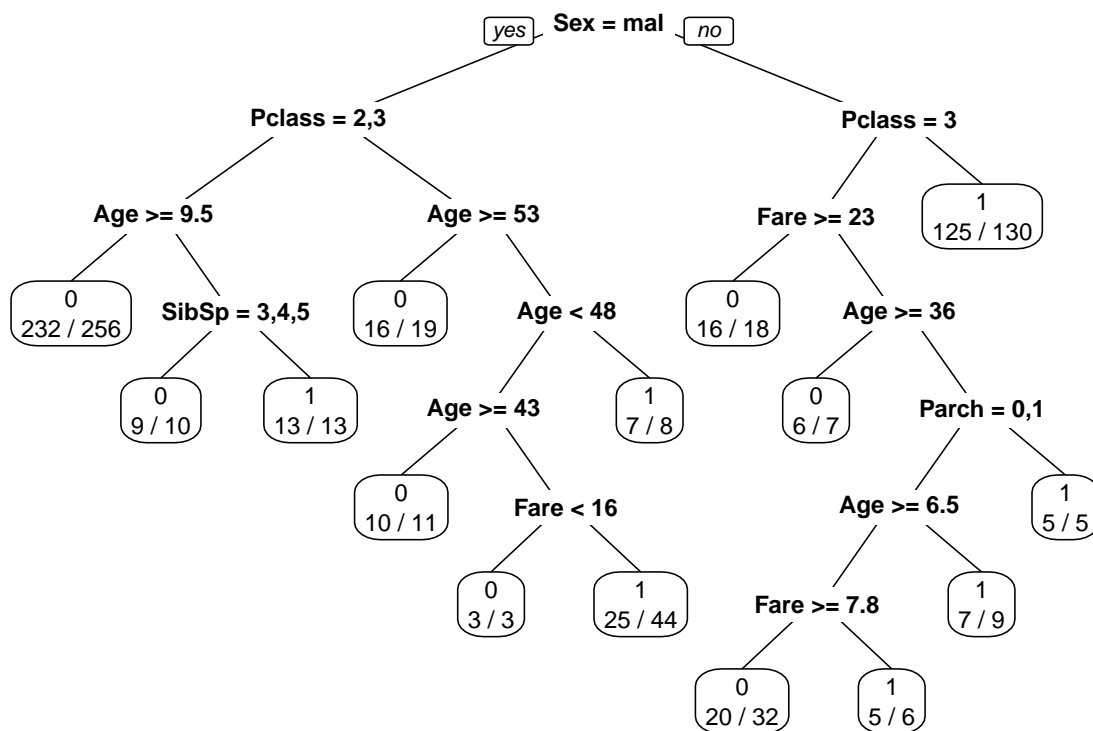
```
random_tree_modified=rpart(train_set,cost = c(3,1,1,2,1,1,1),control = rpart.control(minsplit = 10, minb
                        maxcompete = 4, maxsurrogate = 5, usesurrogate = 2, xval = 10,
                        surrogatestyle = 0, maxdepth = 10))
prp(random_tree_modified,extra=2)
```

Question 5: Try different combinations of values for some of the parameters (decreasing minsplit, minbucket, cp and cost values, for example) and check the performance of each combination on the testing subset. How does this performance change? How do the obtained trees change? Is there any relationship between the parameters values and the shape of the trees?


```
##      Balanced Accuracy : 0.7265
##
##      'Positive' Class : 0
##
```

The accuracy obtained is similar than the obtained with the default values. We picked a small value for cp, so the tree has a lot of branches and it is very deep.

```
random_tree_modified=rpart(train_set,cost = c(1,1,1,1,1,1),control = rpart.control(minsplit = 10, minbu
maxcompete = 4, maxsurrogate = 5, usesurrogate = 2, xval = 10,
surrogatestyle = 0, maxdepth = 10))
prp(random_tree_modified,extra=2)
```



```
testPredMod <- predict(random_tree_modified, test_set,type="class")
confusionMatrix(testPredMod,test_set$Survived)
```

```
## Confusion Matrix and Statistics
##
##      Reference
## Prediction  0  1
##      0  71 19
##      1 13 38
##
##      Accuracy : 0.773
##      95% CI : (0.695, 0.8393)
```



```

##      No Information Rate : 0.5957
##      P-Value [Acc > NIR] : 6.723e-06
##
##              Kappa : 0.5207
##  McNemar's Test P-Value : 0.3768
##
##      Sensitivity : 0.8452
##      Specificity : 0.6667
##      Pos Pred Value : 0.7889
##      Neg Pred Value : 0.7451
##      Prevalence : 0.5957
##      Detection Rate : 0.5035
##      Detection Prevalence : 0.6383
##      Balanced Accuracy : 0.7560
##
##      'Positive' Class : 0
##

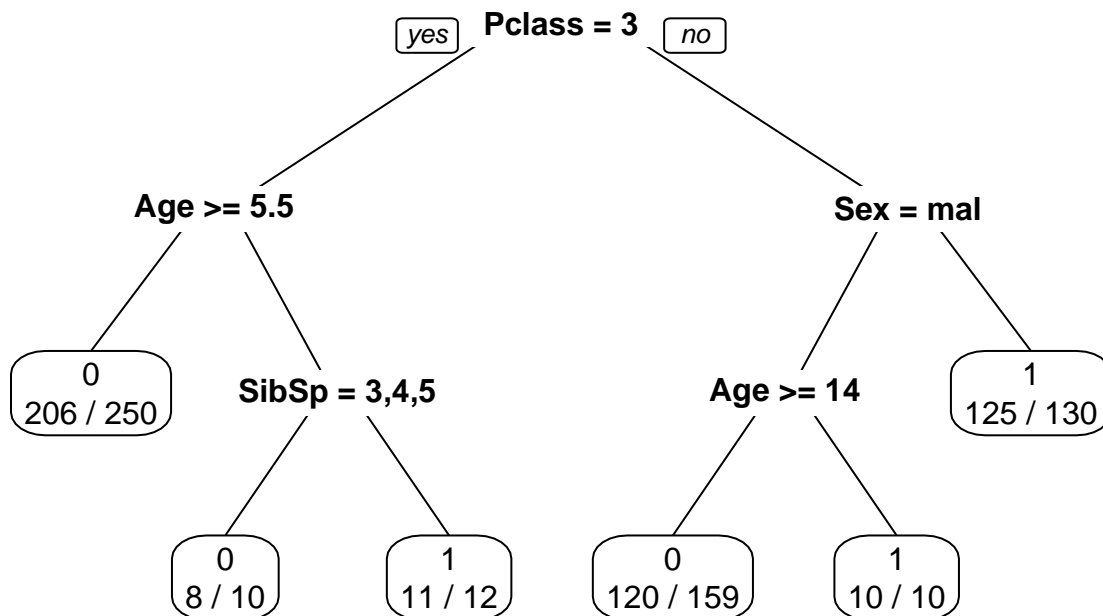
```

In this other case the cp value is not so low, so the overfitting is reduced. That's why we have a better accuracy (around 77%), which is higher than the one obtained with the default parameters.

```

random_tree_modified=rpart(train_set,cost = c(1,3,1,1,1,1),control = rpart.control(minsplit = 10, minbu
maxcompete = 4, maxsurrogate = 5, usesurrogate = 2, xval = 10,
surrogatestyle = 0, maxdepth = 10))
prp(random_tree_modified,extra=2)

```




```
testPredMod <- predict(random_tree_modified, test_set,type="class")
confusionMatrix(testPredMod,test_set$Survived)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0  1
##           0 80 29
##           1  4 28
##
##           Accuracy : 0.766
##           95% CI : (0.6873, 0.8331)
##    No Information Rate : 0.5957
##    P-Value [Acc > NIR] : 1.551e-05
##
##           Kappa : 0.4772
##  McNemar's Test P-Value : 2.943e-05
##
##           Sensitivity : 0.9524
##           Specificity : 0.4912
##           Pos Pred Value : 0.7339
##           Neg Pred Value : 0.8750
##           Prevalence : 0.5957
##           Detection Rate : 0.5674
##    Detection Prevalence : 0.7730
##           Balanced Accuracy : 0.7218
##
##           'Positive' Class : 0
##
```

Here the accuracy is a little bit worse. This might be because we've forced the Pclass to be more important than the sex assigning to it a higher cost. We also can see a big difference between the sensitivity and the specificity. That's because deaths are predicted worse than the survivals.

```
random_tree_modified=rpart(train_set,cost = c(1,3,1,1,1,1),control = rpart.control(minsplit = 40, minbu
maxcompete = 4, maxsurrogate = 5, usesurrogate = 2, xval = 10,
surrogatestyle = 0, maxdepth = 10))
prp(random_tree_modified,extra=2)
```



```
testPredMod <- predict(random_tree_modified, test_set,type="class")
confusionMatrix(testPredMod,test_set$Survived)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0  1
##           0 80 34
##           1  4 23
##
##           Accuracy : 0.7305
##           95% CI : (0.6493, 0.8017)
##       No Information Rate : 0.5957
##       P-Value [Acc > NIR] : 0.000584
##
##           Kappa : 0.3888
##  McNemar's Test P-Value : 2.546e-06
##
##           Sensitivity : 0.9524
##           Specificity : 0.4035
##       Pos Pred Value : 0.7018
##       Neg Pred Value : 0.8519
##           Prevalence : 0.5957
##       Detection Rate : 0.5674
##       Detection Prevalence : 0.8085
```

```
##      Balanced Accuracy : 0.6779
##
##      'Positive' Class : 0
##
```

In this last example we have a high value of `cp`, so our tree is underfitted. That's why the accuracy value is worse.

Summarizing, when we have a small `cp` the trees are deeper. The trees also change depending on the costs assigned, so if we assign a high cost value to a variable, the probability of splitting the tree according to that variable would be higher.

Question 6: Which are the combinations of parameters values tested by the `train()` function? Are there any changes in the performance of the algorithm when different combinations of values are used (according to the results of the cross validation)? We change the `tuneLength` to have more parameters tested by default.

```
fitControl <- trainControl(## 10-fold CV
                           method = "repeatedcv",
                           number = 10,
                           ## repeated ten times
                           repeats = 1)
treeFit <- train(train_set[,2:7],train_set[,1],
                 method = "rpartCost",
                 trControl = fitControl,tuneLength = 8)
treeFit
```

```
## Cost-Sensitive CART
##
## 571 samples
## 6 predictor
## 2 classes: '0', '1'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold, repeated 1 times)
## Summary of sample sizes: 514, 513, 514, 514, 514, 514, ...
## Resampling results across tuning parameters:
##
## Cost cp Accuracy Kappa Accuracy SD Kappa SD
## 1 0.00000000 0.7931942 0.5678883 0.049218917 0.10406017
## 1 0.06617192 0.7722928 0.5139953 0.051797259 0.10405088
## 1 0.13234385 0.7828191 0.5419774 0.062604256 0.13021897
## 1 0.19851577 0.7828191 0.5419774 0.062604256 0.13021897
## 1 0.26468769 0.7828191 0.5419774 0.062604256 0.13021897
## 1 0.33085962 0.7828191 0.5419774 0.062604256 0.13021897
## 1 0.39703154 0.7828191 0.5419774 0.062604256 0.13021897
## 2 0.00000000 0.8019056 0.5721818 0.066182059 0.14518498
## 2 0.06617192 0.8055051 0.5648747 0.035941035 0.08519026
## 2 0.13234385 0.8055051 0.5648747 0.035941035 0.08519026
## 2 0.19851577 0.8055051 0.5648747 0.035941035 0.08519026
## 2 0.26468769 0.6305505 0.1210261 0.058084662 0.19578535
## 2 0.33085962 0.5954628 0.0000000 0.003252191 0.00000000
## 2 0.39703154 0.5954628 0.0000000 0.003252191 0.00000000
```

```
## 3 0.00000000 0.8106776 0.5845858 0.045037441 0.10349900
## 3 0.06617192 0.8055051 0.5648747 0.035941035 0.08519026
## 3 0.13234385 0.8055051 0.5648747 0.035941035 0.08519026
## 3 0.19851577 0.8055051 0.5648747 0.035941035 0.08519026
## 3 0.26468769 0.5954628 0.0000000 0.003252191 0.00000000
## 3 0.33085962 0.5954628 0.0000000 0.003252191 0.00000000
## 3 0.39703154 0.5954628 0.0000000 0.003252191 0.00000000
## 4 0.00000000 0.8124622 0.5849283 0.035583772 0.08304251
## 4 0.06617192 0.8055051 0.5648747 0.035941035 0.08519026
## 4 0.13234385 0.8055051 0.5648747 0.035941035 0.08519026
## 4 0.19851577 0.8055051 0.5648747 0.035941035 0.08519026
## 4 0.26468769 0.5954628 0.0000000 0.003252191 0.00000000
## 4 0.33085962 0.5954628 0.0000000 0.003252191 0.00000000
## 4 0.39703154 0.5954628 0.0000000 0.003252191 0.00000000
## 5 0.00000000 0.8159407 0.5931636 0.039210661 0.09075601
## 5 0.06617192 0.8055051 0.5648747 0.035941035 0.08519026
## 5 0.13234385 0.8055051 0.5648747 0.035941035 0.08519026
## 5 0.19851577 0.8055051 0.5648747 0.035941035 0.08519026
## 5 0.26468769 0.5954628 0.0000000 0.003252191 0.00000000
## 5 0.33085962 0.5954628 0.0000000 0.003252191 0.00000000
## 5 0.39703154 0.5954628 0.0000000 0.003252191 0.00000000
## 6 0.00000000 0.8037810 0.5619537 0.035227988 0.08511210
## 6 0.06617192 0.8055051 0.5648747 0.035941035 0.08519026
## 6 0.13234385 0.8055051 0.5648747 0.035941035 0.08519026
## 6 0.19851577 0.7077435 0.3081931 0.099966261 0.26942824
## 6 0.26468769 0.5954628 0.0000000 0.003252191 0.00000000
## 6 0.33085962 0.5954628 0.0000000 0.003252191 0.00000000
## 6 0.39703154 0.5954628 0.0000000 0.003252191 0.00000000
## 7 0.00000000 0.7827284 0.5098137 0.033902820 0.08267277
## 7 0.06617192 0.8055051 0.5648747 0.035941035 0.08519026
## 7 0.13234385 0.8055051 0.5648747 0.035941035 0.08519026
## 7 0.19851577 0.6673926 0.1994096 0.095380130 0.26095745
## 7 0.26468769 0.5954628 0.0000000 0.003252191 0.00000000
## 7 0.33085962 0.5954628 0.0000000 0.003252191 0.00000000
## 7 0.39703154 0.5954628 0.0000000 0.003252191 0.00000000
## 8 0.00000000 0.7844828 0.5124330 0.037388036 0.09146136
## 8 0.06617192 0.8055051 0.5648747 0.035941035 0.08519026
## 8 0.13234385 0.8055051 0.5648747 0.035941035 0.08519026
## 8 0.19851577 0.6270417 0.0892509 0.067699908 0.18950992
## 8 0.26468769 0.5954628 0.0000000 0.003252191 0.00000000
## 8 0.33085962 0.5954628 0.0000000 0.003252191 0.00000000
## 8 0.39703154 0.5954628 0.0000000 0.003252191 0.00000000
##
```

```
## Accuracy was used to select the optimal model using the largest value.
```

```
## The final values used for the model were cp = 0 and Cost = 5.
```

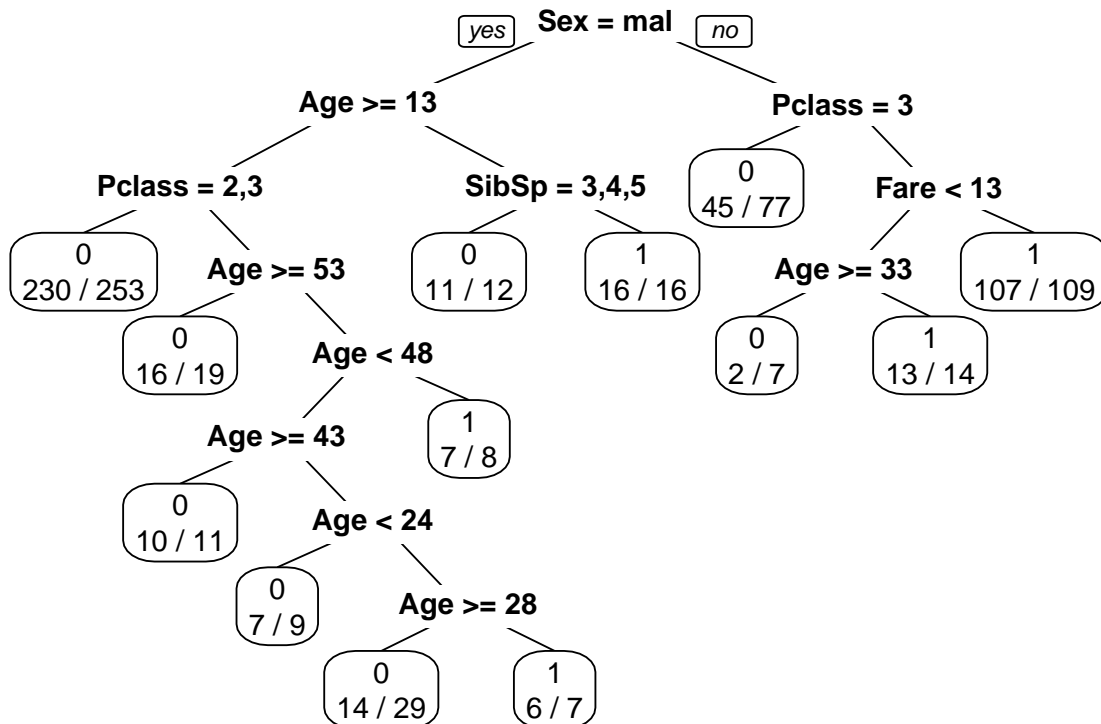
When different combinations are used the performance changes, as we can see above. The combination with better performance is the one picked for the final model.

```
treeFit$bestTune
```

Question 7: Which is the final combination of parameters values used? Which is the shape of the tree trained with this automatic tuning function? (CLUE: if the result of calling to the `train()` function is stored in an object called `rpartFit`, check the `rpartFit$finalModel` object).

```
##      cp Cost
## 29  0    5
```

```
prp(treeFit$finalModel,extra=2)
```



```
testPredTreeFit <- predict(treeFit$finalModel, test_set,type="class")
confusionMatrix(testPredTreeFit,test_set$Survived)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0  1
##           0 75 32
##           1  9 25
##
##           Accuracy : 0.7092
##           95% CI : (0.6268, 0.7826)
##           No Information Rate : 0.5957
##           P-Value [Acc > NIR] : 0.0034120
##
##           Kappa : 0.3544
```

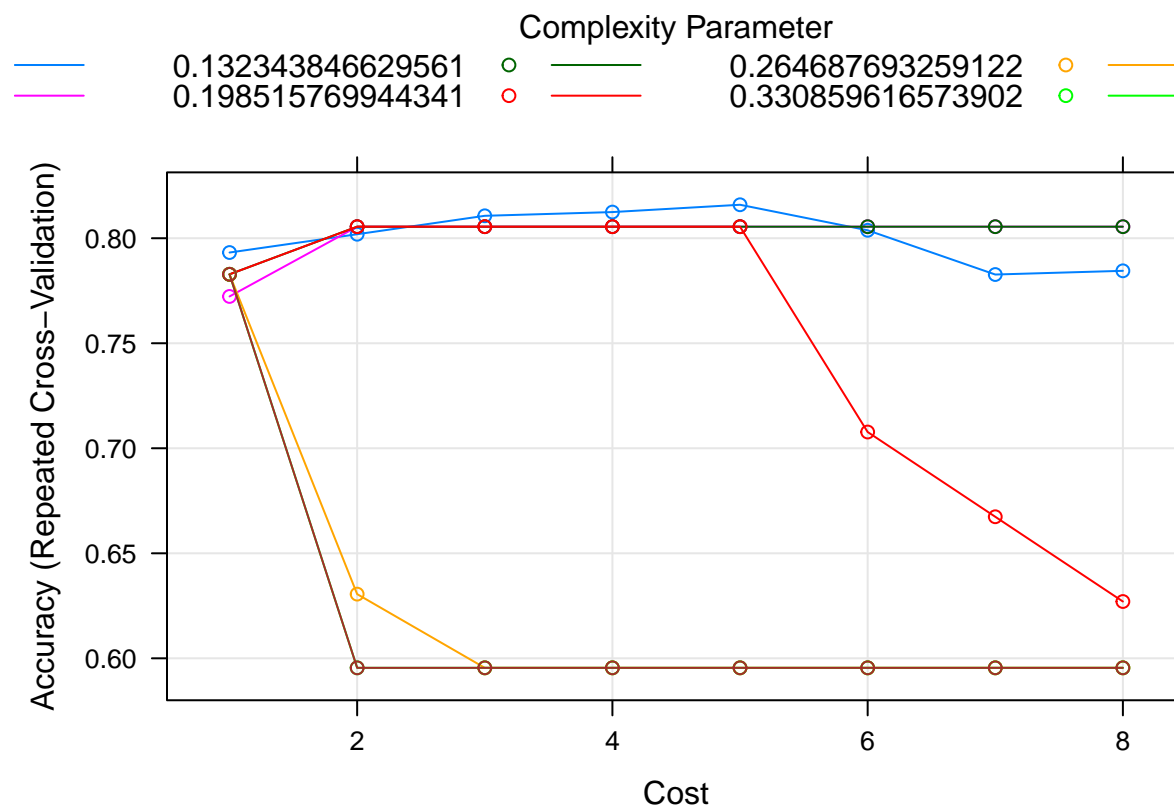
```
## McNemar's Test P-Value : 0.0005908
##
##      Sensitivity : 0.8929
##      Specificity : 0.4386
##      Pos Pred Value : 0.7009
##      Neg Pred Value : 0.7353
##      Prevalence : 0.5957
##      Detection Rate : 0.5319
##      Detection Prevalence : 0.7589
##      Balanced Accuracy : 0.6657
##
##      'Positive' Class : 0
##
```

The final values used for the model were $cp = 0$ and $Cost = 5$ because the accuracy there is the highest.

We checked the accuracy in our testing data set and it is around 70%. The main variables used are sex, age and Pclass. We can see the shape of the tree obtained in the image above.

```
plot(treeFit)
```

Question 8: Plot the result of calling the `train()` function with the `plot()` function. What does



this plot represent?

We can see here the variation of the accuracy for the different combinations used in the train function. The

one picked is the one with better accuracy and less complexity (in case several combinations have the same accuracy values).

Question 9: Answer Question 6 again but now with the results of this new run. (tuning the parameters)

Question 10: Analogously to Question 9, respond to Question 7 with the results of this new run.

```
treeGrid <- expand.grid(Cost=c(1, 2, 3, 5, 10),
cp = c(0, 0.01, 0.02, 0.04, 0.07, 0.10))
T1<-Sys.time()
treeFit2 <- train(train_set[,2:7],train_set[,1],
                 method = "rpartCost",
                 trControl = fitControl,
                 tuneGrid = treeGrid)
T2<-Sys.time()
timeDecisionTree=difftime(T2,T1)
treeFit2
```

Question 11: Answer Question 8 again but now with the results of this new run.

```
## Cost-Sensitive CART
##
## 571 samples
## 6 predictor
## 2 classes: '0', '1'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold, repeated 1 times)
## Summary of sample sizes: 513, 514, 514, 514, 514, 514, ...
## Resampling results across tuning parameters:
##
## Cost cp Accuracy Kappa Accuracy SD Kappa SD
## 1 0.00 0.8054749 0.5953473 0.05857563 0.12065422
## 1 0.01 0.7897762 0.5556960 0.04704474 0.10162248
## 1 0.02 0.8038113 0.5679461 0.02984207 0.06750979
## 1 0.04 0.7932849 0.5440617 0.04227041 0.08344708
## 1 0.07 0.7828191 0.5423179 0.04982924 0.10015239
## 2 0.00 0.7967937 0.5667897 0.04444794 0.08636616
## 2 0.01 0.8090744 0.5796051 0.02818998 0.06576674
## 2 0.02 0.8143376 0.5871147 0.03340694 0.07754291
## 2 0.04 0.8055656 0.5655358 0.02823179 0.06616199
## 2 0.07 0.8055656 0.5655358 0.02823179 0.06616199
## 3 0.00 0.8125832 0.5918952 0.04637065 0.10022165
## 3 0.01 0.8143376 0.5876819 0.03340694 0.07778032
## 3 0.02 0.8160920 0.5907112 0.03344046 0.07759869
## 3 0.04 0.8055656 0.5655358 0.02823179 0.06616199
## 3 0.07 0.8055656 0.5655358 0.02823179 0.06616199
## 5 0.00 0.8143376 0.5889401 0.03969463 0.09088321
```



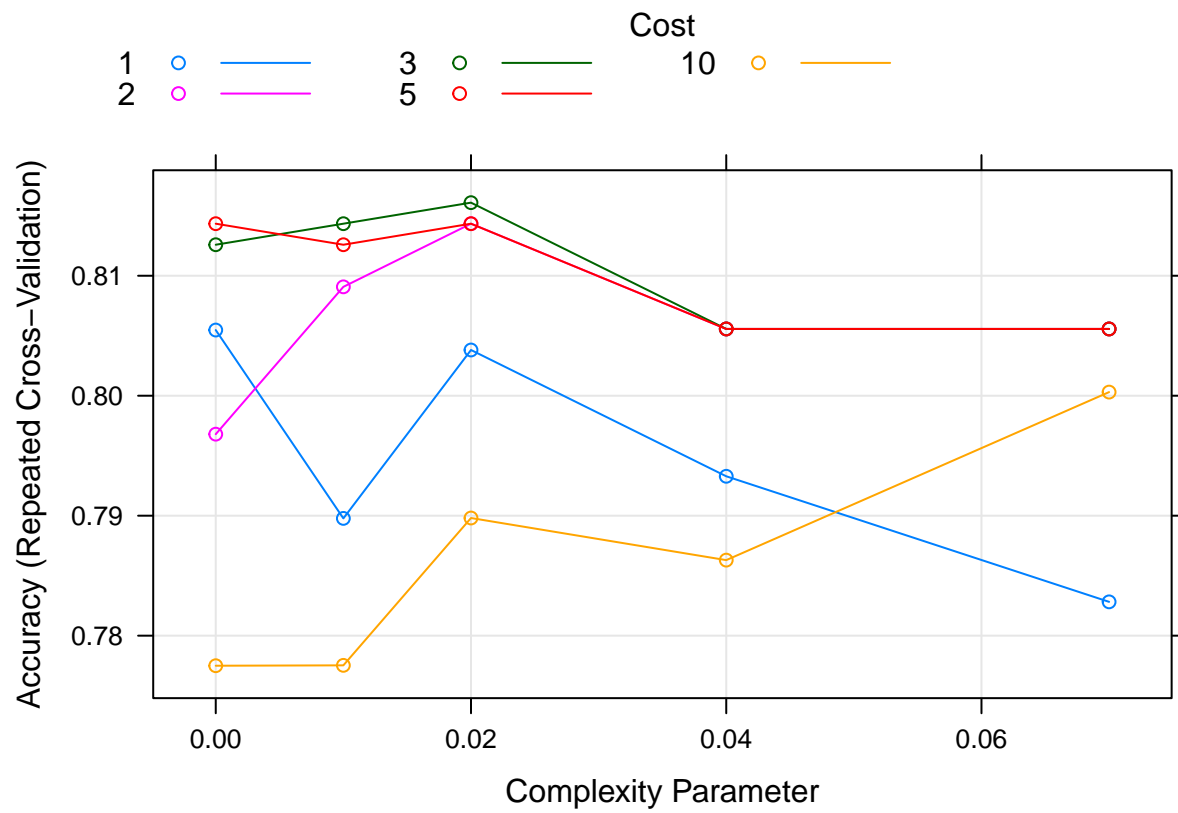
```
##      5      0.01  0.8125832  0.5820159  0.03622340  0.08523496
##      5      0.02  0.8143376  0.5863302  0.03539517  0.08294011
##      5      0.04  0.8055656  0.5655358  0.02823179  0.06616199
##      5      0.07  0.8055656  0.5655358  0.02823179  0.06616199
##     10      0.00  0.7774955  0.4952884  0.04865154  0.11877539
##     10      0.01  0.7775257  0.4962626  0.03734329  0.09040739
##     10      0.02  0.7898064  0.5263073  0.02768478  0.06659218
##     10      0.04  0.7862976  0.5184680  0.02487103  0.06028931
##     10      0.07  0.8003025  0.5525916  0.02918680  0.06951318
##
## Accuracy was used to select the optimal model using the largest value.
## The final values used for the model were cp = 0.02 and Cost = 3.
```

```
testPredTreeFit2 <- predict(treeFit2$finalModel, test_set,type="class")
confusionMatrix(testPredTreeFit2,test_set$Survived)
```

```
## Confusion Matrix and Statistics
##
##              Reference
## Prediction  0  1
##           0 79 30
##           1  5 27
##
##              Accuracy : 0.7518
##              95% CI : (0.6721, 0.8206)
##      No Information Rate : 0.5957
##      P-Value [Acc > NIR] : 7.372e-05
##
##              Kappa : 0.4456
##      McNemar's Test P-Value : 4.976e-05
##
##              Sensitivity : 0.9405
##              Specificity : 0.4737
##              Pos Pred Value : 0.7248
##              Neg Pred Value : 0.8438
##              Prevalence : 0.5957
##              Detection Rate : 0.5603
##      Detection Prevalence : 0.7730
##              Balanced Accuracy : 0.7071
##
##      'Positive' Class : 0
##
```

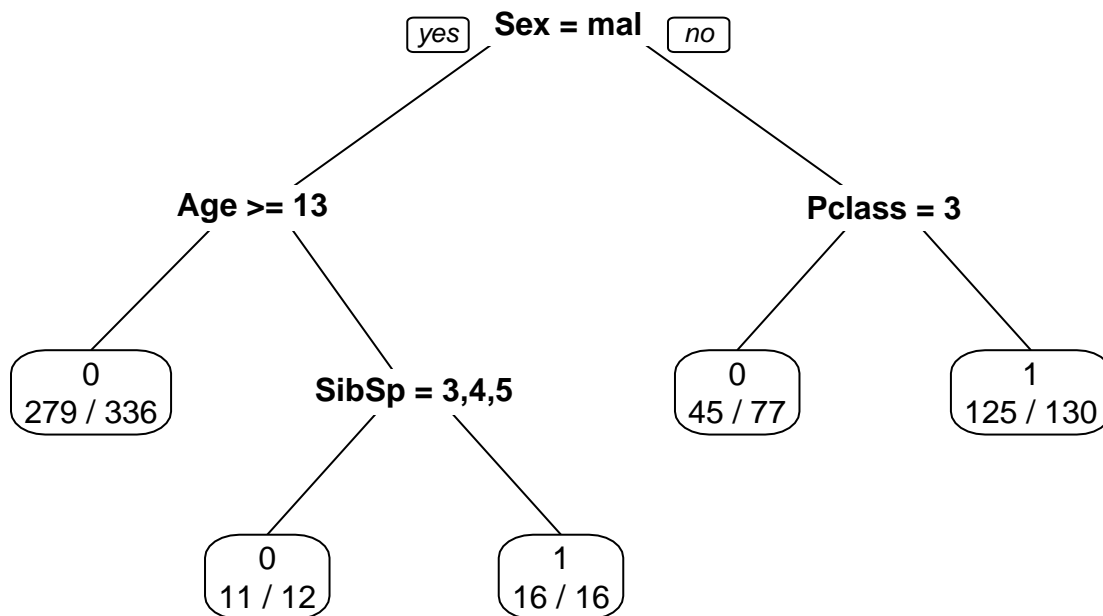
The accuracy of the best combination is around 81%. We can also check the accuracy value in our testing set, which is better than the one from the previous question.

```
plot(treeFit2)
```



We can see the combination picked above and the plot of the different combination, where we can see which one was the best.

```
prp(treeFit2$finalModel,extra=2)
```



The shape of the tree is plotted above. The tree obtained is simple and we can see that the rule of being a woman and not being in third class is still present. We can also see that male children with 3,4 or 5 siblings died, but the ones with less siblings survived.

Training a random forest

Question 12: Answer Question 6 again but now with the results of this new run.

Question 13: Which is the final combination of the parameters used?

Question 14: Answer Question 8 again but now with the results of this new run. We load the random forest library.

```
library(randomForest)
```

```
fitControl <- trainControl(## 10-fold CV
                           method = "repeatedcv",
                           number = 10,
                           ## repeated ten times
                           repeats = 1)

T1<-Sys.time()
treeFitRandom <- train(train_set[,2:7],train_set[,1],
                      method = "rf",
                      trControl = fitControl,
```

```

        ntree=2000,
        tuneLength = 5)
T2<-Sys.time()
timeRandomForest=difftime(T2,T1)

treeFitRandom

```

```

## Random Forest
##
## 571 samples
## 6 predictor
## 2 classes: '0', '1'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold, repeated 1 times)
## Summary of sample sizes: 514, 514, 514, 514, 514, 513, ...
## Resampling results across tuning parameters:
##
## mtry Accuracy Kappa Accuracy SD Kappa SD
## 2 0.8265275 0.6293173 0.05718243 0.12101680
## 3 0.8265578 0.6327745 0.04194302 0.08688276
## 4 0.8265880 0.6344535 0.03266540 0.06631840
## 5 0.8178463 0.6173292 0.03815299 0.07832908
## 6 0.8108590 0.6039120 0.03859495 0.07861150
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 4.

```

Above we can see the value for mtry picked in the best combination (the one with highest accuracy).

```

testPredTreeFitRandom <- predict(treeFitRandom$finalModel, test_set,type="class")
confusionMatrix(testPredTreeFitRandom,test_set$Survived)

```

```


## Confusion Matrix and Statistics
##
##           Reference
## Prediction 0 1
##           0 67 18
##           1 17 39
##
##           Accuracy : 0.7518
##           95% CI : (0.6721, 0.8206)
##           No Information Rate : 0.5957
##           P-Value [Acc > NIR] : 7.372e-05
##
##           Kappa : 0.4832
##           McNemar's Test P-Value : 1
##
##           Sensitivity : 0.7976
##           Specificity : 0.6842
##           Pos Pred Value : 0.7882
##           Neg Pred Value : 0.6964
##           Prevalence : 0.5957

```

```
##      Detection Rate : 0.4752
##      Detection Prevalence : 0.6028
##      Balanced Accuracy : 0.7409
##
##      'Positive' Class : 0
##
```

The accuracy is around 75%.

```
plot(treeFitRandom)
```



Here we can see the combinations plotted and the best one is clear in the plot.

```
treeRandomForest=getTree(treeFitRandom$finalModel,1,labelVar = TRUE)
treeRandomForest
```

```
##      left daughter right daughter split var split point status prediction
## 1          2          3      Pclass      3.00000      1      <NA>
## 2          4          5       Fare     16.89790      1      <NA>
## 3          6          7       Age      1.50000      1      <NA>
## 4          8          9      Parch      1.00000      1      <NA>
## 5         10         11       Sex      1.00000      1      <NA>
## 6          0          0      <NA>      0.00000     -1       1
## 7         12         13       Sex      1.00000      1      <NA>
## 8         14         15       Fare     13.75000      1      <NA>
```

## 9	0	0	<NA>	0.00000	-1	1
## 10	16	17	Age	2.50000	1	<NA>
## 11	18	19	Age	18.00000	1	<NA>
## 12	20	21	Fare	20.80000	1	<NA>
## 13	22	23	Age	32.50000	1	<NA>
## 14	24	25	Sex	1.00000	1	<NA>
## 15	0	0	<NA>	0.00000	-1	0
## 16	0	0	<NA>	0.00000	-1	0
## 17	26	27	Fare	26.12500	1	<NA>
## 18	0	0	<NA>	0.00000	-1	1
## 19	28	29	Pclass	1.00000	1	<NA>
## 20	30	31	Fare	16.40000	1	<NA>
## 21	32	33	Parch	32.00000	1	<NA>
## 22	34	35	Age	30.50000	1	<NA>
## 23	0	0	<NA>	0.00000	-1	0
## 24	36	37	Age	36.00000	1	<NA>
## 25	38	39	Fare	12.93750	1	<NA>
## 26	40	41	Age	43.00000	1	<NA>
## 27	0	0	<NA>	0.00000	-1	1
## 28	42	43	Fare	127.81665	1	<NA>
## 29	44	45	SibSp	2.00000	1	<NA>
## 30	46	47	Fare	13.43750	1	<NA>
## 31	48	49	SibSp	3.00000	1	<NA>
## 32	50	51	Fare	30.25625	1	<NA>
## 33	0	0	<NA>	0.00000	-1	0
## 34	52	53	Age	24.50000	1	<NA>
## 35	54	55	Fare	7.83750	1	<NA>
## 36	0	0	<NA>	0.00000	-1	1
## 37	56	57	Age	41.50000	1	<NA>
## 38	58	59	Age	21.00000	1	<NA>
## 39	60	61	Fare	13.25000	1	<NA>
## 40	0	0	<NA>	0.00000	-1	1
## 41	62	63	Parch	1.00000	1	<NA>
## 42	64	65	Age	39.00000	1	<NA>
## 43	0	0	<NA>	0.00000	-1	0
## 44	66	67	Fare	26.12500	1	<NA>
## 45	0	0	<NA>	0.00000	-1	0
## 46	68	69	Parch	1.00000	1	<NA>
## 47	70	71	Fare	15.67500	1	<NA>
## 48	0	0	<NA>	0.00000	-1	1
## 49	0	0	<NA>	0.00000	-1	0
## 50	0	0	<NA>	0.00000	-1	0
## 51	0	0	<NA>	0.00000	-1	1
## 52	72	73	Age	9.50000	1	<NA>
## 53	74	75	Age	27.50000	1	<NA>
## 54	0	0	<NA>	0.00000	-1	0
## 55	76	77	Fare	8.20625	1	<NA>
## 56	0	0	<NA>	0.00000	-1	0
## 57	0	0	<NA>	0.00000	-1	1
## 58	78	79	Age	18.50000	1	<NA>
## 59	0	0	<NA>	0.00000	-1	0
## 60	80	81	Age	30.50000	1	<NA>
## 61	0	0	<NA>	0.00000	-1	0
## 62	0	0	<NA>	0.00000	-1	0

## 63	0	0	<NA>	0.00000	-1	1
## 64	82	83	Fare	64.97915	1	<NA>
## 65	84	85	Fare	73.86460	1	<NA>
## 66	86	87	Age	33.00000	1	<NA>
## 67	0	0	<NA>	0.00000	-1	0
## 68	88	89	Fare	10.87920	1	<NA>
## 69	0	0	<NA>	0.00000	-1	1
## 70	0	0	<NA>	0.00000	-1	0
## 71	90	91	Fare	15.97500	1	<NA>
## 72	92	93	Fare	24.82500	1	<NA>
## 73	94	95	Fare	7.22710	1	<NA>
## 74	96	97	Age	26.50000	1	<NA>
## 75	98	99	Fare	8.77500	1	<NA>
## 76	0	0	<NA>	0.00000	-1	1
## 77	0	0	<NA>	0.00000	-1	0
## 78	0	0	<NA>	0.00000	-1	0
## 79	0	0	<NA>	0.00000	-1	1
## 80	0	0	<NA>	0.00000	-1	0
## 81	100	101	Age	41.00000	1	<NA>
## 82	102	103	Fare	27.06875	1	<NA>
## 83	104	105	Fare	86.08540	1	<NA>
## 84	106	107	Age	75.50000	1	<NA>
## 85	108	109	Fare	81.33750	1	<NA>
## 86	0	0	<NA>	0.00000	-1	1
## 87	0	0	<NA>	0.00000	-1	0
## 88	110	111	Fare	7.76250	1	<NA>
## 89	0	0	<NA>	0.00000	-1	1
## 90	0	0	<NA>	0.00000	-1	1
## 91	0	0	<NA>	0.00000	-1	0
## 92	0	0	<NA>	0.00000	-1	1
## 93	112	113	Fare	30.25625	1	<NA>
## 94	114	115	Age	20.50000	1	<NA>
## 95	116	117	Age	19.50000	1	<NA>
## 96	118	119	Fare	7.83540	1	<NA>
## 97	0	0	<NA>	0.00000	-1	1
## 98	0	0	<NA>	0.00000	-1	0
## 99	120	121	Age	28.75000	1	<NA>
## 100	122	123	Age	32.50000	1	<NA>
## 101	0	0	<NA>	0.00000	-1	0
## 102	0	0	<NA>	0.00000	-1	1
## 103	124	125	Fare	28.87500	1	<NA>
## 104	126	127	Age	27.50000	1	<NA>
## 105	0	0	<NA>	0.00000	-1	1
## 106	128	129	Age	57.00000	1	<NA>
## 107	0	0	<NA>	0.00000	-1	1
## 108	130	131	SibSp	1.00000	1	<NA>
## 109	132	133	Fare	101.18960	1	<NA>
## 110	134	135	Fare	6.98750	1	<NA>
## 111	136	137	Age	50.00000	1	<NA>
## 112	0	0	<NA>	0.00000	-1	0
## 113	138	139	Age	2.50000	1	<NA>
## 114	0	0	<NA>	0.00000	-1	0
## 115	140	141	Fare	7.13335	1	<NA>
## 116	142	143	Parch	1.00000	1	<NA>

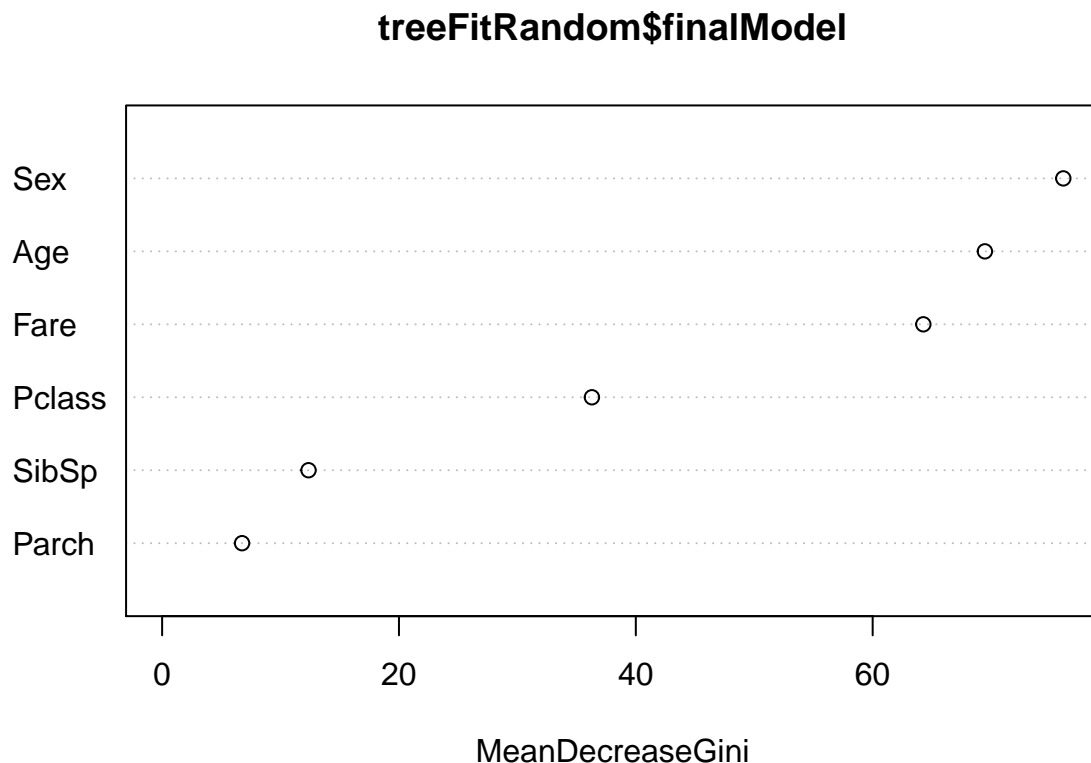
## 117	0	0	<NA>	0.00000	-1	0
## 118	144	145	Fare	7.75835	1	<NA>
## 119	146	147	SibSp	1.00000	1	<NA>
## 120	0	0	<NA>	0.00000	-1	0
## 121	0	0	<NA>	0.00000	-1	1
## 122	0	0	<NA>	0.00000	-1	1
## 123	0	0	<NA>	0.00000	-1	1
## 124	0	0	<NA>	0.00000	-1	0
## 125	148	149	Age	34.00000	1	<NA>
## 126	150	151	Fare	77.00835	1	<NA>
## 127	0	0	<NA>	0.00000	-1	0
## 128	152	153	Fare	52.27710	1	<NA>
## 129	0	0	<NA>	0.00000	-1	0
## 130	0	0	<NA>	0.00000	-1	0
## 131	0	0	<NA>	0.00000	-1	1
## 132	154	155	Age	47.00000	1	<NA>
## 133	0	0	<NA>	0.00000	-1	0
## 134	0	0	<NA>	0.00000	-1	0
## 135	0	0	<NA>	0.00000	-1	1
## 136	156	157	Age	16.00000	1	<NA>
## 137	0	0	<NA>	0.00000	-1	1
## 138	0	0	<NA>	0.00000	-1	0
## 139	158	159	Age	6.00000	1	<NA>
## 140	0	0	<NA>	0.00000	-1	0
## 141	0	0	<NA>	0.00000	-1	1
## 142	160	161	Age	18.50000	1	<NA>
## 143	0	0	<NA>	0.00000	-1	0
## 144	0	0	<NA>	0.00000	-1	0
## 145	0	0	<NA>	0.00000	-1	1
## 146	162	163	Age	25.50000	1	<NA>
## 147	0	0	<NA>	0.00000	-1	0
## 148	0	0	<NA>	0.00000	-1	1
## 149	164	165	Fare	52.82710	1	<NA>
## 150	0	0	<NA>	0.00000	-1	1
## 151	0	0	<NA>	0.00000	-1	0
## 152	166	167	Fare	43.68125	1	<NA>
## 153	168	169	Fare	59.05210	1	<NA>
## 154	0	0	<NA>	0.00000	-1	0
## 155	0	0	<NA>	0.00000	-1	1
## 156	0	0	<NA>	0.00000	-1	0
## 157	170	171	Age	20.00000	1	<NA>
## 158	0	0	<NA>	0.00000	-1	1
## 159	0	0	<NA>	0.00000	-1	0
## 160	0	0	<NA>	0.00000	-1	0
## 161	172	173	Fare	8.10415	1	<NA>
## 162	0	0	<NA>	0.00000	-1	0
## 163	174	175	Fare	13.34165	1	<NA>
## 164	0	0	<NA>	0.00000	-1	1
## 165	0	0	<NA>	0.00000	-1	0
## 166	176	177	Age	49.00000	1	<NA>
## 167	0	0	<NA>	0.00000	-1	0
## 168	0	0	<NA>	0.00000	-1	1
## 169	0	0	<NA>	0.00000	-1	0
## 170	178	179	Fare	7.81460	1	<NA>

## 171	180	181	Fare	8.29375	1	<NA>
## 172	182	183	Fare	7.85000	1	<NA>
## 173	0	0	<NA>	0.00000	-1	0
## 174	0	0	<NA>	0.00000	-1	0
## 175	0	0	<NA>	0.00000	-1	1
## 176	184	185	Age	46.00000	1	<NA>
## 177	186	187	Age	53.50000	1	<NA>
## 178	0	0	<NA>	0.00000	-1	0
## 179	0	0	<NA>	0.00000	-1	1
## 180	0	0	<NA>	0.00000	-1	1
## 181	0	0	<NA>	0.00000	-1	0
## 182	0	0	<NA>	0.00000	-1	0
## 183	0	0	<NA>	0.00000	-1	1
## 184	188	189	Fare	29.36040	1	<NA>
## 185	0	0	<NA>	0.00000	-1	0
## 186	0	0	<NA>	0.00000	-1	1
## 187	190	191	Fare	33.09790	1	<NA>
## 188	192	193	Age	42.50000	1	<NA>
## 189	194	195	Age	42.50000	1	<NA>
## 190	0	0	<NA>	0.00000	-1	0
## 191	0	0	<NA>	0.00000	-1	1
## 192	0	0	<NA>	0.00000	-1	0
## 193	0	0	<NA>	0.00000	-1	0
## 194	0	0	<NA>	0.00000	-1	1
## 195	0	0	<NA>	0.00000	-1	0

Random forests are a black box, so we can't visualize the output. If we want to have an idea of how the final tree would look like we can use the `getTree` function and get a representation of a sample tree from the random forest (as done above).

```
varImpPlot(treeFitRandom$finalModel)
```

Question 15: Plot the importance of each variable for the model with function `VarImpPlot()` from package `caret`. Which are the most relevant variables according to their mean decrease of the Gini index? Are these variables the ones selected when we built our decision trees?



Sex is the most important parameter, while Parch and SibSp are the least important ones. The most important ones were the ones selected in the trees before. There's a big gap between the most important variables and the less important ones and that is linked with the depth where they appear in the trees.

Question 16: Answer Question 6 again but now with the results of this new run. (tuning the parameters)

Question 17: Answer Question 13 again but now with the results of this new run.

Question 18: Answer Question 8 again but now with the results of this new run. The mtry is the number of variables per level, so we are trying to tune it to see when we have the best results.


```
treeGridRF <- expand.grid(mtry=c(2,3,4,5,6))
fitControl <- trainControl(## 10-fold CV
                           method = "repeatedcv",
                           number = 10,
                           ## repeated ten times
                           repeats = 1)

set.seed(100)
treeFitRandom2 <- train(train_set[,2:7],train_set[,1],
                       method = "rf",
                       trControl = fitControl,
                       ntree=2000,
                       tuneGrid = treeGridRF)

treeFitRandom2
```

```
## Random Forest
##
## 571 samples
## 6 predictor
## 2 classes: '0', '1'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold, repeated 1 times)
## Summary of sample sizes: 514, 514, 514, 514, 514, 514, ...
## Resampling results across tuning parameters:
##
## mtry Accuracy Kappa Accuracy SD Kappa SD
## 2 0.8232305 0.6234820 0.05393756 0.1176481
## 3 0.8337871 0.6482701 0.06290706 0.1354272
## 4 0.8284936 0.6394213 0.06449337 0.1371823
## 5 0.8319722 0.6467077 0.05325205 0.1146701
## 6 0.8319722 0.6460280 0.04638761 0.1012548
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 3.
```


```
plot(treeFitRandom2)
```



```
treeRandomForest2=getTree(treeFitRandom2$finalModel,1,labelVar = TRUE)
```

We've observed that if we train it in a different way we have different results. This is weird, as we think if we set the seed the results should be the same.

```
set.seed(100)
treeFitRandomTEST <- train(as.factor(Survived)~.,data=train_set[,1:7],
  method = "rf",
  trControl = fitControl,
  ntree=2000,
  tuneGrid = treeGridRF)
plot(treeFitRandomTEST)
```



```
testPredTreeFitRandom2 <- predict(treeFitRandom2$finalModel, test_set,type="class")
confusionMatrix(testPredTreeFitRandom2,test_set$Survived)
```

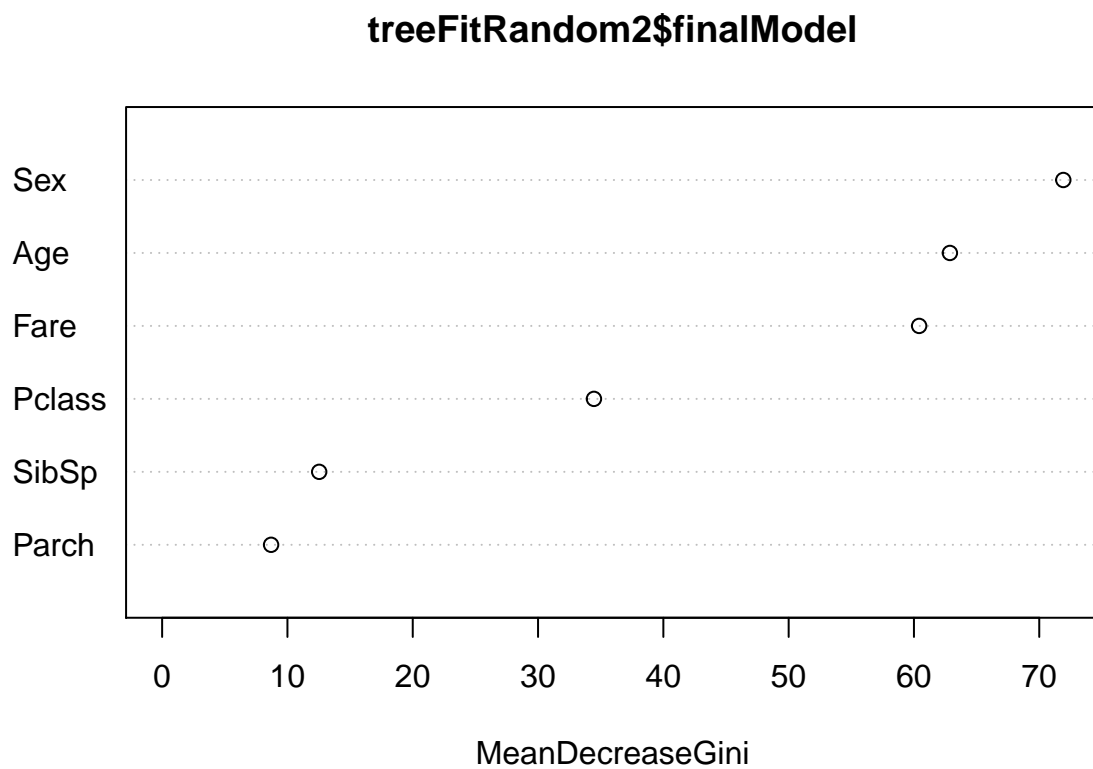
```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0  1
##           0 70 20
##           1 14 37
##
##           Accuracy : 0.7589
##           95% CI : (0.6797, 0.8269)
##           No Information Rate : 0.5957
##           P-Value [Acc > NIR] : 3.444e-05
```

```
##
##           Kappa : 0.4908
## Mcnemar's Test P-Value : 0.3912
##
##           Sensitivity : 0.8333
##           Specificity : 0.6491
##           Pos Pred Value : 0.7778
##           Neg Pred Value : 0.7255
##           Prevalence : 0.5957
##           Detection Rate : 0.4965
##           Detection Prevalence : 0.6383
##           Balanced Accuracy : 0.7412
##
##           'Positive' Class : 0
##
```

The accuracy have increased slightly compared with the run we did before.

```
varImpPlot(treeFitRandom2$finalModel)
```

Question 19: Answer Question 15 again but now with the results of this new run.



The importance of the variables is almost the same. Age and fare are closer and sex a little bit further, but in general the importances are the same.

Question 20: Which is the difference in performance with regards to the testing subset between the best decision tree model and the best random forest model? With decision trees we got the best accuracy in the second tree we did in question 4 (with an accuracy of 0.773). With random forest we achieved less accuracy (0.7589). We think this is because the dataset is too simple for using random forest. Even if there was a slight improvement, still decision trees would be better in this case because of the time it takes to run.

The time it tooks in the cases we got the best results with each method are:

```
timeDecisionTree
```

```
## Time difference of 4.105108 secs
```

```
timeRandomForest
```

```
## Time difference of 2.310153 mins
```